

AI3603: Computer Vision, Spring 2025
Indian Institute of Technology Hyderabad
Homework 2, Generative Modelling
25 points. Assigned 16.03.2025, Due **11:59 pm on 22.03.2025**

The essence of the scientific spirit is to realize what a wonderful world it is that we live in. – Sir C. V. Raman

Instructions:

- It is **strongly recommended** that you work on your homework on an *individual* basis. If you have any questions or concerns, feel free to talk to the instructor or the TAs.
- Use the MNIST dataset.
- You are free to use Copilot. Please turn in your prompts.
- Please turn in Python Notebooks with the following notation for the file name: `your-roll-number-hw2.ipynb`.

Questions:

1. Variational Autoencoder (VAE) (10)

- (a) Train a VAE on the MNIST dataset. Describe your architecture. Use the reparameterization trick to ensure gradient propagation.
- (b) We discussed in class that $q(z|x)$ is a Gaussian distribution that whose parameters are learned to be “close” to the standard normal distribution. Show that the latent distribution $q(z|x)$ does indeed vary with the conditioning variable x . Why is this important?
- (c) Visualise the generative process of the VAE by sampling from the latent space. Do “nearby” samples in the latent space correspond to “similar” samples in the image space?
- (d) Verify qualitatively and quantitatively that the model has indeed learned to generate *diverse samples of good quality*.

2. Denoising Diffusion Probabilistic Model (DDPM) (10)

- (a) Train a DDPM on the MNIST dataset. Use a U-Net model and Gaussian noise for diffusion.
- (b) Experiment with different time steps for the diffusion process. How does the number of steps affect the quality and diversity of the generated samples?
- (c) Visualise the generative process of the DDPM by starting from a noise image. Use an appropriate number of steps for the diffusion process.
- (d) Do “nearby” samples in the latent space lead to “similar” samples in the image space?
- (e) Verify qualitatively and quantitatively that the model has indeed learned to generate *diverse samples of good quality*.

3. Comparing DDPM and VAE (5)

- (a) Fundamentally, how does a DDPM differ from a VAE?
- (b) How does a DDPM compare with a VAE in terms of sample quality and diversity?
- (c) How does a DDPM compare with a VAE in terms of training time?